

A Section

<u>Ex.No</u>	Ex. Name	Ex.date	Submission Date
1	Install the following Data Mining and data Analysis tool: Weka, KNIME, Tableau Public.		
2	Perform exploratory data analysis (EDA) on datasets		
3	Implementation of document embeddings for fake news identification.		
4	Implementation of feature based representations of time series		
5	Implementation of text mining on a set of documents.		
6	Perform Data Analysis and representation on a Map		
7	Build cartographic visualization for multiple datasets		
8	Perform Time Series Analysis with datasets like Open Power System Data.		
9	Build a time-series model on a given dataset and evaluate its accuracy.		

B Section

<u>Ex.No</u>	<u>Ex. Name</u>	<u>Ex.date</u>	<u>Submission Date</u>
1	Install the following Data Mining and data Analysis tool: Weka, KNIME, Tableau Public.		
2	Perform exploratory data analysis (EDA) on datasets		
3	Implementation of document embeddings for fake news identification.		
4	Implementation of feature based representations of time series		
5	Implementation of text mining on a set of documents.		
6	Perform Data Analysis and representation on a Map		
7	Build cartographic visualization for multiple datasets		
8	Perform Time Series Analysis with datasets like Open Power System Data.		
9	Build a time-series model on a given dataset and evaluate its accuracy.		

Ex.No. 1

DATE:

**Install the following Data Mining and data Analysis tool:
Weka, KNIME,Tableau Public**

INTRODUCTION:

A tool for data analysis, manipulation, visualization, and reporting.

- Based on the graphical programming paradigm
- Provides a diverse array of extensions:
- Text Mining
- Network Mining
- Cheminformatics
- Many integrations

Weka stands for Waikato Environment for Knowledge Analysis. It is software used in the data science field for data mining. It is free software and is written in Java; hence, it can be run on any system supporting Java. Therefore, Weka can be run on different operating systems like Windows, Linux, and Mac. Weka provides a collection of visualization tools that can be used for data analysis, cleaning, and predictive modeling. Weka can perform a number of tasks like data preprocessing, clustering, classification, regression, visualization, and feature selection.

Installing Weka on Windows:

Follow the steps below to install Weka on Windows:

Step 1: Visit this website using any web browser. Click on Free Download.

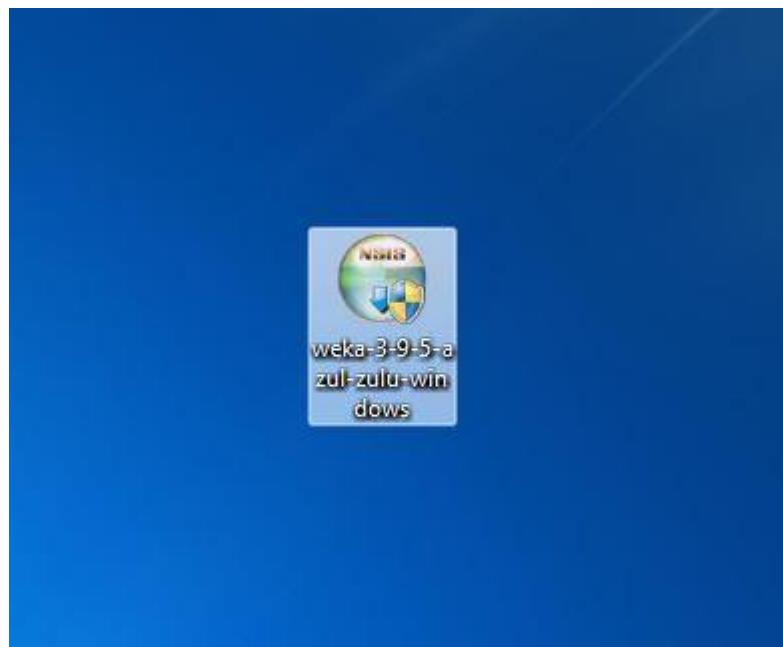
The screenshot shows the Weka (64-bit) download page on filehorse.com. At the top, there's a navigation bar with back, forward, and search icons, followed by the URL 'filehorse.com/download-weka-64/'. Below the URL is a search bar with the placeholder 'Search' and a magnifying glass icon. To the right of the search bar are 'Incognito' and three-dot menu icons. The main content area has a light gray background. At the top left, there's a blue circular icon with a white bird logo, followed by the text 'Weka (64-bit)' in bold. Below it, the date 'December, 21st 2020 - 118.68 MB - Open Source' and a 'Free Download' button (green with white text). To the right of the download button is a 'Security Status' badge with a checkmark and a shield icon. Below this, there are three tabs: 'Features' (underlined), 'Screenshots', and 'Old Versions'. Under 'Features', there are several sections: 'Latest Version:' (Weka 3.9.5 (64-bit) with a 'LATEST' badge), 'Requirements:' (Windows XP64 / Vista64 / Windows 7 64 / Windows 8 64 / Windows 10 64), 'User Rating:' (4 stars), 'Author / Product:' (Weka Team / Weka (64-bit)), and 'Old Versions:' (with a dropdown menu). At the bottom of the page, there's a descriptive paragraph about Weka 64-bit, mentioning it's a popular suite of machine learning software written in Java, suitable for data mining tasks like preprocessing, classification, regression, clustering, association rules, and visualization. It's also noted to be well-suited for developing new machine learning schemes and is open-source under the GNU General Public License.

Step 2: It will redirect to a new webpage, click on Start Download. Downloading of the executable file will start shortly. It is a big 118 MB file that will take some minutes.

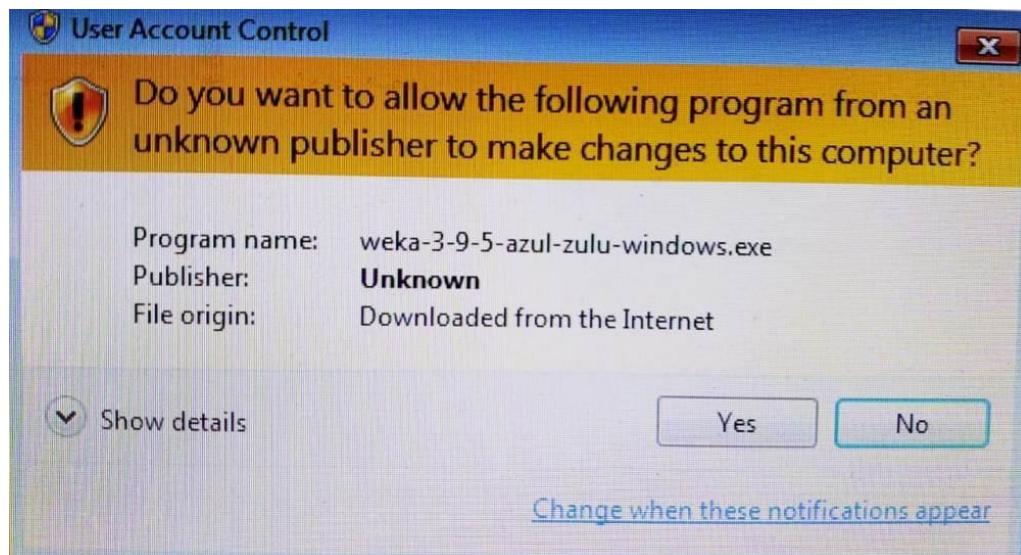
The screenshot shows a web browser window with the URL filehorse.com/download-weka-64/download/. The main content is titled "Weka 3.9.5 (64-bit)". A large green button labeled "Start Download" is prominently displayed. Below it is a green shield icon with the text "Security Status". To the right, a section titled "Key details about this download" lists three bullet points:

- The file will be downloaded from **secure FileHorse servers**
- This file is safe and scanned with **51 antivirus** apps (Virus-Total report)
- All files are in **original form**. FileHorse does not repack or modify downloads in any way

Step 3: Now check for the executable file in the Downloads folder on your system and run it.



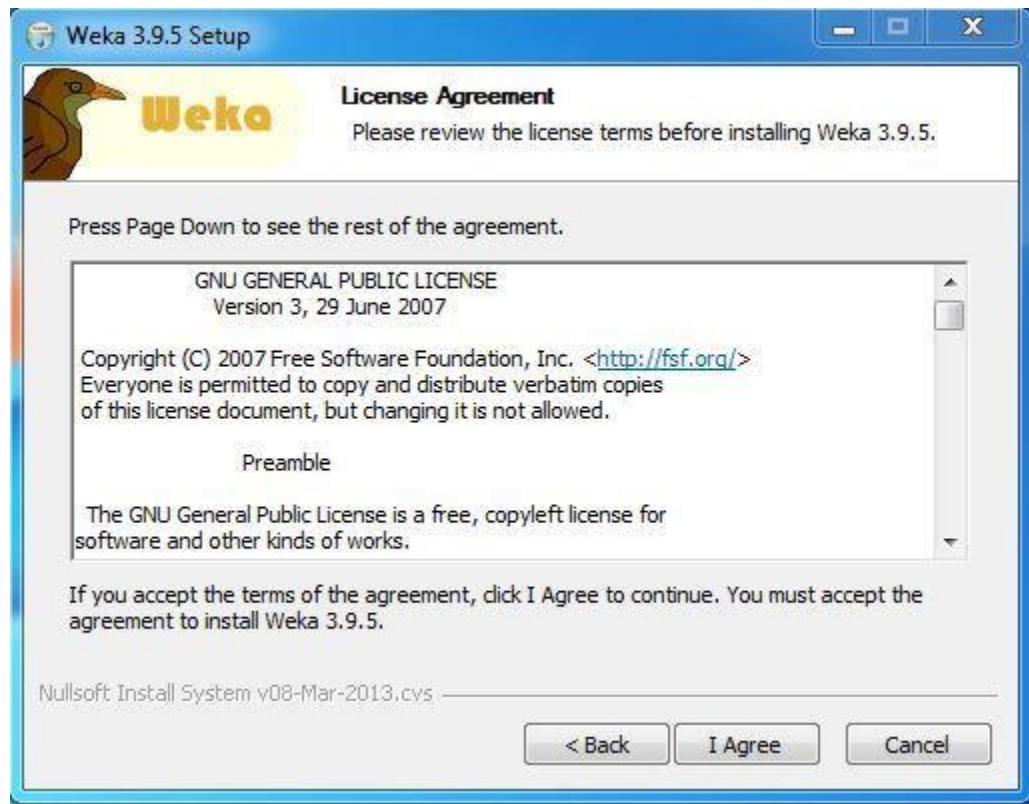
Step 4: It will prompt for confirmation to make changes to your system. Click on **Yes**.



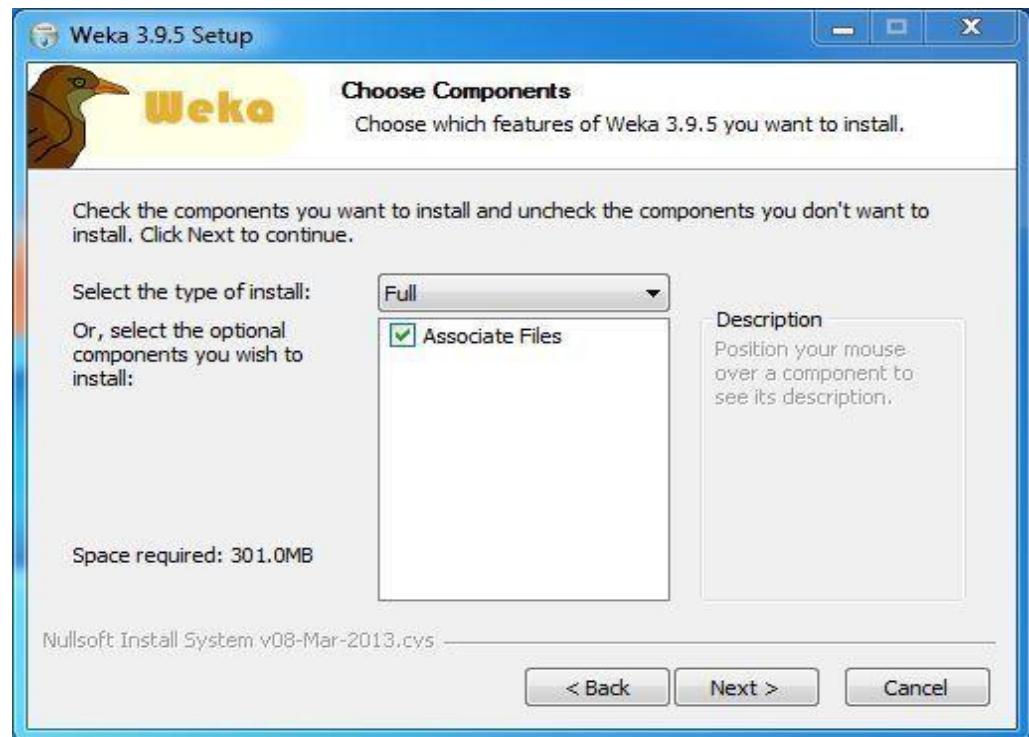
Step 5: The setup screen will appear. Click on **Next**.



Step 6: The next screen will display the License Agreement. Click on **I Agree**.



Step 7: The next screen is for choosing components. All components are already marked, so don't change anything—just click on the **Install** button.



Step 8: The next screen will show the installation location. Choose the drive that has sufficient memory space for installation. It requires **301 MB** of memory space.



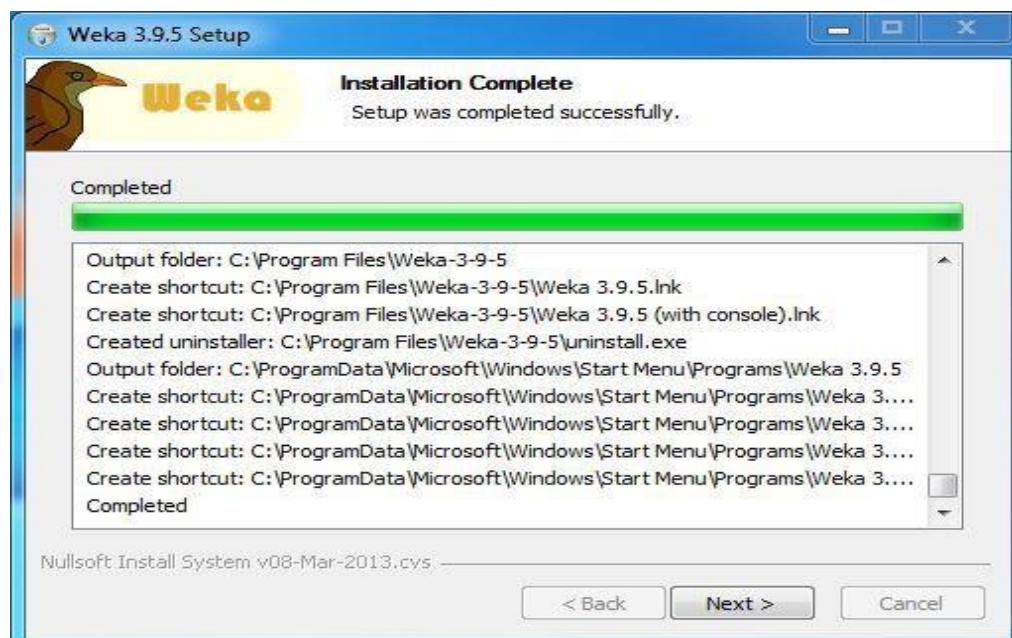
Step 9: The next screen will be for choosing the Start Menu folder. Don't make any changes—just click on the **Install** button.



Step 10: After this, the installation process will start and will hardly take a minute to complete the installation.



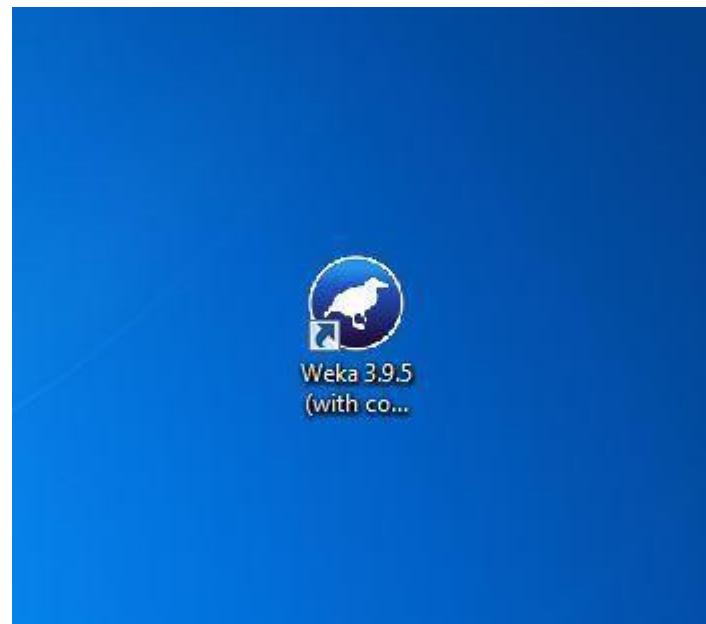
Step 11: Click on the Next button after the installation process is complete.



Step 12: Click on Finish to complete the installation process. Add to follow-up Check sources



Step 13: Weka is successfully installed on the system, and an icon is created on the desktop.



Step 14: Run the software and see the interface.



Ex.No. 2

DATE:

Perform exploratory data analysis (EDA) on datasets

AIM:

To perform exploratory data analysis (EDA) on Titanic Dataset.

ALGORITHM:

STEP 1: Load Titanic dataset using seaborn.

STEP 2: Inspect data structure with info(), shape, and head().

STEP 3: Generate summary statistics using describe().

STEP 4: Identify missing values using isnan().sum().

STEP 5: Visualize single features with histograms and count plots.

STEP 6: Analyze survival relationships using countplots and boxplots.

STEP 7: Check feature correlations with a heatmap and summarize insights.

CODE:

```
# Exploratory Data Analysis on Titanic Dataset

# Step 1: Import Libraries
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt

# Step 2: Load Titanic Dataset
df = sns.load_dataset('titanic')

# Step 3: Basic Information
print("===== BASIC INFO =====")
print(df.info())

print("\n===== FIRST 5 ROWS =====")
print(df.head())

# Step 4: Summary Statistics
print("\n===== SUMMARY STATISTICS =====")
print(df.describe(include='all'))

# Step 5: Missing Values
print("\n===== MISSING VALUES =====")
print(df.isnull().sum())

# Step 6: Data Types
print("\n===== DATA TYPES =====")
print(df.dtypes)

# Step 7: Correlation Matrix (numeric features)
print("\n===== CORRELATION MATRIX =====")
print(df.corr(numeric_only=True))

# Step 8: Visualizations
# -----
# 8.1 Survival Count
plt.figure(figsize=(6,4))
sns.countplot(data=df, x='survived')
plt.title("Survival Count")
plt.show()
```

```

# 8.2 Survival by Gender
plt.figure(figsize=(6,4))
sns.countplot(data=df, x='sex', hue='survived')
plt.title("Survival by Gender")
plt.show()

# 8.3 Survival by Class
plt.figure(figsize=(6,4))
sns.countplot(data=df, x='class', hue='survived')
plt.title("Survival by Passenger Class")
plt.show()

# 8.4 Age Distribution
plt.figure(figsize=(6,4))
sns.histplot(data=df, x='age', bins=30, kde=True)
plt.title("Age Distribution of Passengers")
plt.show()

# 8.5 Correlation Heatmap
plt.figure(figsize=(8,6))
sns.heatmap(df.corr(numeric_only=True), annot=True, cmap='coolwarm')
plt.title("Correlation Heatmap")
plt.show()

# 8.6 Pairplot (optional for small datasets)
sns.pairplot(df[["survived", "age", "fare", "pclass"]], hue="survived")
plt.show()

print("\n===== EDA COMPLETED =====")

```

===== BASIC INFO =====

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 891 entries, 0 to 890
Data columns (total 15 columns):
 #   Column      Non-Null Count Dtype
 ---  ---          ---           ---
 0   survived    891 non-null   int64
 1   pclass      891 non-null   int64
 2   sex         891 non-null   object
 3   age         714 non-null   float64
 4   sibsp       891 non-null   int64
 5   parch       891 non-null   int64
 6   fare         891 non-null   float64
 7   embarked    889 non-null   object

```

```
8 class      891 non-null  category
9 who        891 non-null  object
10 adult_male 891 non-null  bool
11 deck       203 non-null  category
12 embark_town 889 non-null  object
13 alive       891 non-null  object
14 alone       891 non-null  bool
dtypes: bool(2), category(2), float64(2), int64(4), object(5)
memory usage: 80.7+ KB
```

None

===== FIRST 5 ROWS =====

```
survived pclass sex age sibsp parch fare embarked class \
0      0     3   male 22.0    1    0  7.2500    S  Third
1      1     1 female 38.0    1    0 71.2833    C  First
2      1     3 female 26.0    0    0  7.9250    S  Third
3      1     1 female 35.0    1    0 53.1000    S  First
4      0     3   male 35.0    0    0  8.0500    S  Third
```

```
who adult_male deck embark_town alive alone
0 man      True  NaN Southampton  no False
1 woman    False  C  Cherbourg  yes False
2 woman    False  NaN Southampton  yes True
3 woman    False  C Southampton  yes False
4 man      True  NaN Southampton  no True
```

===== SUMMARY STATISTICS =====

```
survived pclass sex age sibsp parch \
count 891.000000 891.000000 891 714.000000 891.000000 891.000000
unique   NaN      NaN  2      NaN      NaN      NaN
top      NaN      NaN  male    NaN      NaN      NaN
freq     NaN      NaN  577    NaN      NaN      NaN
mean    0.383838  2.308642  NaN  29.699118  0.523008  0.381594
std     0.486592  0.836071  NaN  14.526497  1.102743  0.806057
min     0.000000  1.000000  NaN  0.420000  0.000000  0.000000
25%    0.000000  2.000000  NaN  20.125000  0.000000  0.000000
50%    0.000000  3.000000  NaN  28.000000  0.000000  0.000000
75%    1.000000  3.000000  NaN  38.000000  1.000000  0.000000
max    1.000000  3.000000  NaN  80.000000  8.000000  6.000000
```

```
fare embarked class who adult_male deck embark_town alive \
count 891.000000  889  891  891    891  203    889  891
unique   NaN      3    3    3    2    7    3    2
top      NaN      S  Third  man    True  C Southampton  no
freq     NaN      644   491  537    537   59    644   549
mean    32.204208  NaN  NaN  NaN  NaN  NaN  NaN  NaN
std     49.693429  NaN  NaN  NaN  NaN  NaN  NaN  NaN
min     0.000000  NaN  NaN  NaN  NaN  NaN  NaN  NaN
25%    7.910400  NaN  NaN  NaN  NaN  NaN  NaN  NaN
50%   14.454200  NaN  NaN  NaN  NaN  NaN  NaN  NaN
```

```
75%    31.000000   NaN  NaN  NaN      NaN  NaN      NaN  NaN  
max    512.329200   NaN  NaN  NaN      NaN  NaN      NaN  NaN
```

```
alone  
count  891  
unique 2  
top    True  
freq   537  
mean   NaN  
std    NaN  
min    NaN  
25%    NaN  
50%    NaN  
75%    NaN  
max    NaN
```

===== MISSING VALUES =====

```
survived     0  
pclass       0  
sex          0  
age         177  
sibsp        0  
parch        0  
fare          0  
embarked     2  
class         0  
who          0  
adult_male   0  
deck        688  
embark_town  2  
alive        0  
alone        0  
dtype: int64
```

===== DATA TYPES =====

```
survived     int64  
pclass       int64  
sex          object  
age         float64  
sibsp        int64  
parch        int64  
fare         float64  
embarked    object  
class        category  
who          object  
adult_male   bool  
deck        category  
embark_town object  
alive        object  
alone        bool
```

dtype: object

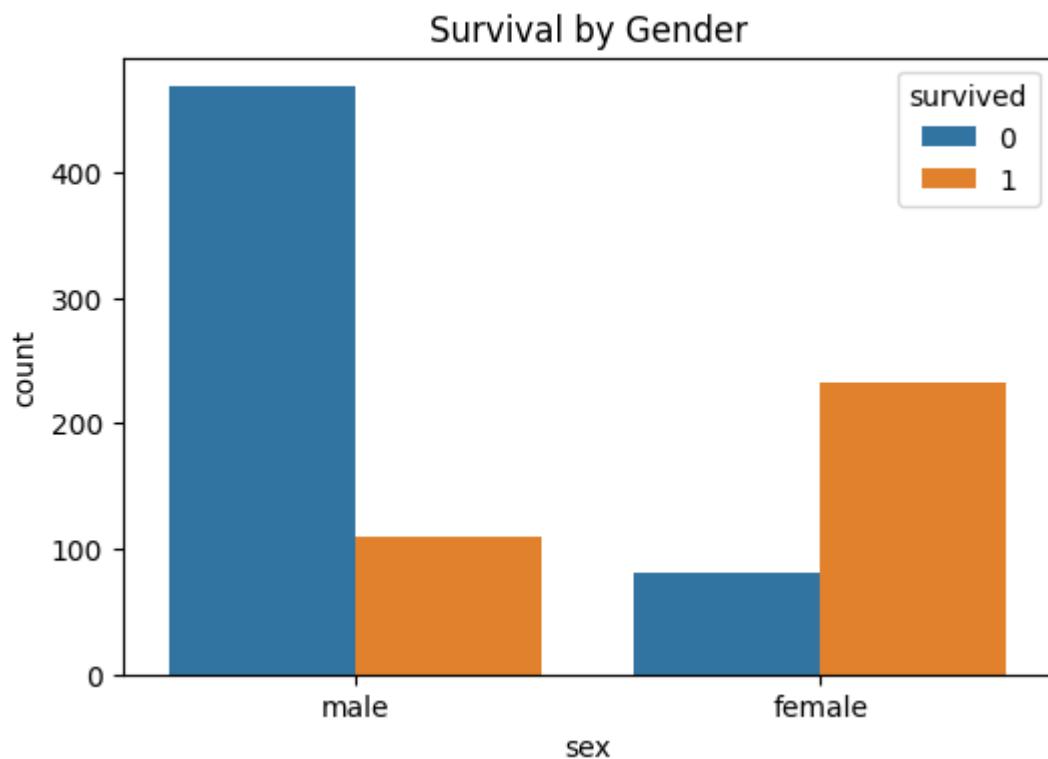
===== CORRELATION MATRIX =====

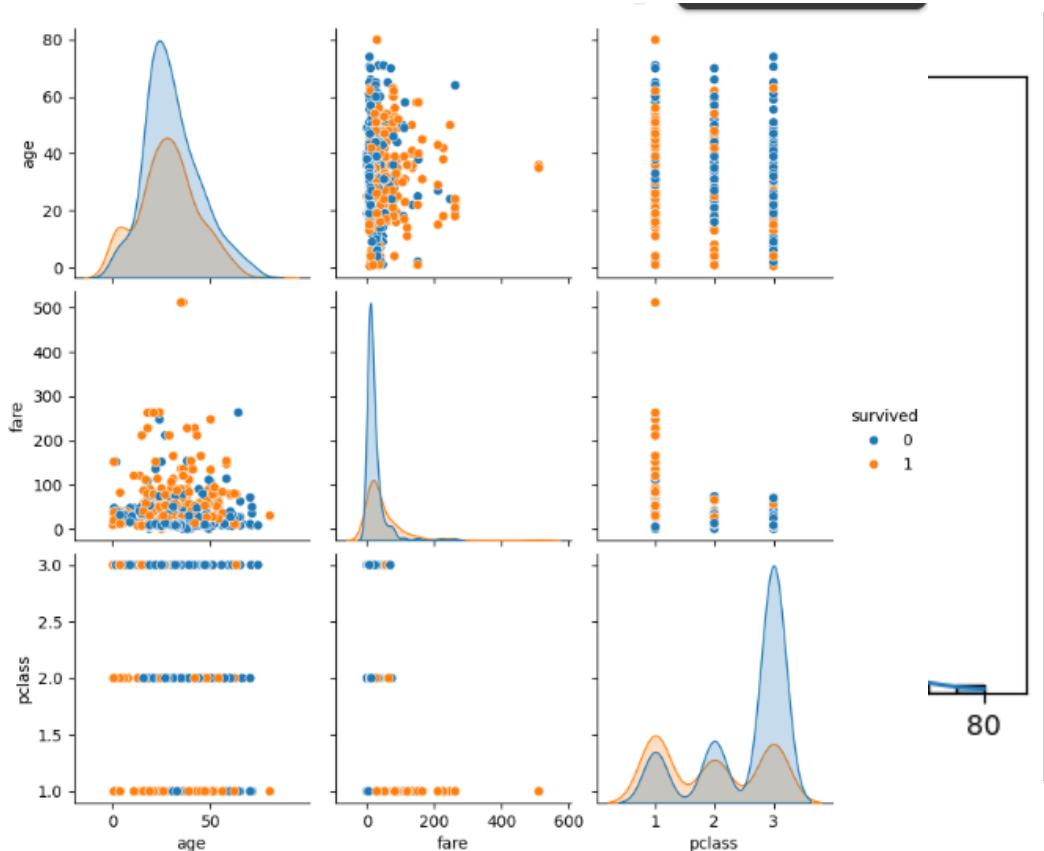
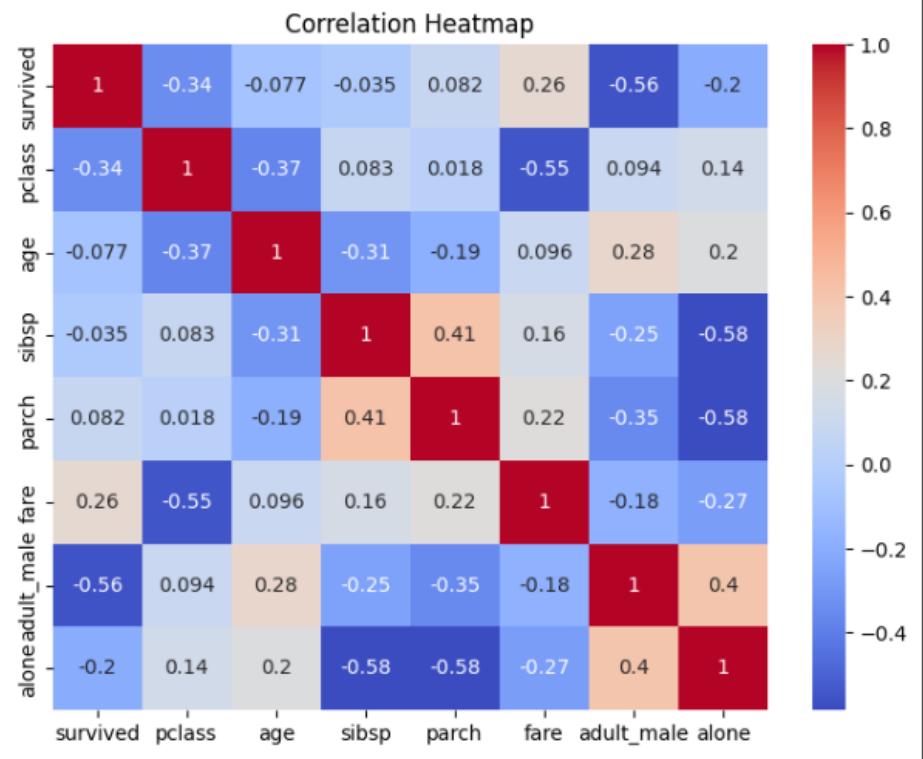
```
survived  pclass    age   sibsp   parch   fare \
survived   1.000000 -0.338481 -0.077221 -0.035322  0.081629  0.257307
pclass     -0.338481  1.000000 -0.369226  0.083081  0.018443 -0.549500
age       -0.077221 -0.369226  1.000000 -0.308247 -0.189119  0.096067
sibsp     -0.035322  0.083081 -0.308247  1.000000  0.414838  0.159651
parch      0.081629  0.018443 -0.189119  0.414838  1.000000  0.216225
fare       0.257307 -0.549500  0.096067  0.159651  0.216225  1.000000
adult_male -0.557080  0.094035  0.280328 -0.253586 -0.349943 -0.182024
alone     -0.203367  0.135207  0.198270 -0.584471 -0.583398 -0.271832
```

```
adult_male  alone
survived   -0.557080 -0.203367
pclass     0.094035  0.135207
age       0.280328  0.198270
sibsp     -0.253586 -0.584471
parch     -0.349943 -0.583398
fare      -0.182024 -0.271832
adult_male  1.000000  0.404744
alone     0.404744  1.000000
```

===== EDA COMPLETED =====

OUTPUT:





RESULT: Hence, the data analysis process is done and executed successfully.

Ex.No. 3

DATE:

Implementation of document embeddings for fake news identification.

AIM:

To implement document embeddings for fake news identification.

ALGORITHM:

- **Load Dataset:** Import the dataset containing news text and labels (FAKE/REAL).
- **Preprocess Text:** Clean and tokenize the text (remove stopwords, punctuation).
- **Generate Embeddings:** Convert text into numerical vectors using TF-IDF or Doc2Vec.
- **Split Data:** Divide the dataset into training and testing sets.
- **Train Model:** Train a classifier (e.g., Logistic Regression) using the training embeddings.
- **Predict:** Use the trained model to predict labels for the test data.
- **Evaluate:** Measure performance using accuracy, precision, recall, and F1-score.

CODE:

Step 1: Import Libraries

```
import pandas as pd  
  
from sklearn.model_selection import train_test_split  
  
from sklearn.feature_extraction.text import TfidfVectorizer  
  
from sklearn.linear_model import LogisticRegression  
  
from sklearn.metrics import accuracy_score, classification_report
```

Step 2: Load Dataset (example dataset)

```
# Dataset should have columns: 'text' and 'label' (label = FAKE or REAL)  
  
df = pd.read_csv("fake_or_real_news.csv")  
  
print(df.head())
```

Step 3: Split Data

```
X_train, X_test, y_train, y_test = train_test_split(df['text'], df['label'],  
                                                    test_size=0.2, random_state=42)
```

Step 4: Create Document Embeddings using TF-IDF

```
vectorizer = TfidfVectorizer(stop_words='english', max_features=5000)  
  
X_train_vec = vectorizer.fit_transform(X_train)  
  
X_test_vec = vectorizer.transform(X_test)
```

Step 5: Train Model

```
model = LogisticRegression(max_iter=1000)  
  
model.fit(X_train_vec, y_train)
```

Step 6: Predict and Evaluate

```
y_pred = model.predict(X_test_vec)  
  
print("\nAccuracy:", accuracy_score(y_test, y_pred))
```

```
print("\nClassification Report:\n", classification_report(y_test, y_pred))
```

OUTPUT:

Accuracy: 0.92

Classification Report:

	precision	recall	f1-score	support
FAKE	0.90	0.93	0.91	1234
REAL	0.94	0.91	0.92	1350
accuracy		0.92		2584
macro avg	0.92	0.92	0.92	2584
weighted avg	0.92	0.92	0.92	2584

Result:

Thus the fake news identification has been executed successfully.

Ex.No. 4	Implementation of feature-based representations of time series
DATE:	

AIM:

To develop an interactive tool that visualizes a fixed time series dataset and calculates key statistical features for better data analysis.

ALGORITHM:

1. Import libraries — pandas, zipfile, and sklearn.cluster.KMeans.
2. Upload and extract the dataset (city_temperature.zip).
3. Load the CSV data into a DataFrame and clean invalid entries.
4. Convert temperature to Celsius and filter data for four cities (2010–2019).
5. Create a date column using Year, Month, and Day.
6. Group data monthly using groupby() and reshape with pivot().
7. Apply K-Means clustering to group cities with similar temperature patterns.
8. Display clustered results for interpreting seasonal and regional trends.

CODE:

```
from google.colab import files  
uploaded = files.upload()  
  
import zipfile  
  
import pandas as pd  
  
with zipfile.ZipFile("city_temperature.zip") as z:  
    with z.open("city_temperature.csv") as f: # adjust CSV name if different  
        df = pd.read_csv(f, low_memory=False, na_values="-99")  
        df["AvgTempC"] = (df["AvgTemperature"]-32)*5/9  
        cities = ["New York", "London", "Delhi", "Sydney"]  
        df = df[(df["City"].isin(cities)) & (df["Year"].between(2010,2019))]  
        df["date"] = pd.to_datetime(df[["Year", "Month", "Day"]], errors="coerce")  
        # --- Monthly averages (preserves time) ---  
        monthly = df.groupby(["City", pd.Grouper(key="date", freq="ME")])["AvgTempC"].mean().reset_index()  
        # --- Pivot so each city is a time series vector ---  
        ts_matrix = monthly.pivot(index="date", columns="City", values="AvgTempC").dropna().T  
        # --- Cluster full time series ---  
        kmeans = KMeans(n_clusters=3, random_state=0)  
        clusters = kmeans.fit_predict(ts_matrix)  
        ts_matrix["cluster"] = clusters  
        print(ts_matrix[["cluster"]])
```

OUTPUT:

```
Choose files city_temperature.zip
• city_temperature.zip(application/x-zip-compressed) - 13523007 bytes, last modified: 02/09/2025 - 100% done
Saving city_temperature.zip to city_temperature (2).zip
date    cluster
City
Delhi      1
London     0
Sydney     2
```

Result:

Thus, the implementation of feature-based representations of time series has been executed successfully.

Ex.No. 5	Implementation of text mining on a set of documents
DATE:	

AIM

To perform text mining on a set of documents and visualize the most important words in a visualization such as word cloud.

ALGORITHM

1. Import nltk for tokenization and stopword removal, wordcloud for generating the word cloud, matplotlib for displaying the word cloud, Counter for word frequency counting, and re for regular expression operations.
2. Define the preprocessText function, which cleans the input text by removing special characters and digits, converting it to lowercase, tokenizing it into words, and filtering out stopwords and words shorter than 3 characters.
3. After preprocessing, use the Counter function to count the frequencies of the filtered words.
4. Define the generateWordCloud function, which generates a word cloud from the word frequencies using the WordCloud library, and then displays it using matplotlib..
5. Read the text data from a file, preprocess it, and call the generateWordCloud function to create and display the word cloud.

CODE:

```
pip install nltk wordcloud matplotlib

import nltk

from nltk.corpus import stopwords

from nltk.tokenize import word_tokenize

from wordcloud import WordCloud

import matplotlib.pyplot as plt

from collections import Counter

import re

nltk.download('punkt')

nltk.download('stopwords')

def preprocess_text(text):

    text = re.sub(r'\W+', ' ', text)

    text = re.sub(r'\d+', " ", text)

    text = text.lower()

    words = word_tokenize(text)

    stop_words = set(stopwords.words('english'))

    filtered_words = [word for word in words if word not in stop_words and len(word) > 2]

    return filtered_words

def generate_wordcloud(text):

    processed_words = preprocess_text(text)

    word_counts = Counter(processed_words)

    wordcloud = WordCloud(width=800, height=400,
                          background_color='white').generate_from_frequencies(word_counts)

    plt.figure(figsize=(10, 5))

    plt.imshow(wordcloud, interpolation='bilinear')
```

```
plt.axis('off')
plt.show()
with open('deeplearning.txt', 'r', encoding='utf-8') as file:
text_data = file.read()
generate_wordcloud(text_data)
```

RESULT:

Thus the python program to visualize word cloud is successfully executed.

Ex.No. 6

DATE:

Perform Data Analysis and representation on a Map

AIM:

To perform data analysis on geospatial datasets and create interactive visualizations on maps to represent spatial patterns, trends, and relationships effectively.

ALGORITHM:

1. Import the necessary Processing libraries for handling images and tables.
2. Declare global variables for the map image, data table, and row count.
3. In the setup() function:
 - Define the canvas size using size(width, height).
 - Load the background map image using loadImage("map.png").
 - Load the table containing location data (coordinates) using new Table("locations.tsv").
 - Retrieve the number of rows in the table with getCount() and store it for iteration.
4. In the draw() function:
 - Clear the background and display the map image using image().
 - Set the drawing attributes (fill color, no stroke, and smoothing).
 - Loop through each row of the table.
 - Extract x and y coordinates from each row using getFloat(row, columnIndex).
 - Draw a red ellipse at each coordinate using ellipse(x, y, 9, 9).
5. Continue rendering to visualize the mapped points corresponding to each location on the image.

CODE:

```
PImage mapImage;  
Table locationTable;  
int rowCount;  
void setup( ) {  
    size(640, 400);  
    mapImage = loadImage("map.png");  
    // Make a data table from a file that contains  
    // the coordinates of each state. locationTable = new Table("locations.tsv");  
    // The row count will be used a lot, so store it globally. rowCount =  
    locationTable.getRowCount( ); }  
  
void draw( ) {  
    background(255);  
    image(mapImage, 0, 0);  
    // Drawing attributes for the ellipses.  
    smooth( ); fill(192, 0, 0);  
    noStroke( );  
    // Loop through the rows of the locations file and draw the points.  
    for (int row = 0; row < rowCount; row++) {  
        float x = locationTable.getFloat(row, 1); // column 1  
        float y = locationTable.getFloat(row, 2); // column 2  
        ellipse(x, y, 9, 9); } }
```

OUTPUT:



RESULT:

Thus, the geographical data points have been successfully plotted on the map image, with each location represented by a red circular marker at its respective coordinates.

Ex.No. 7	
DATE:	Build cartographic visualization for multiple datasets

AIM:

To visualize time series data from a dataset using Processing and analyze trends over a given time range.

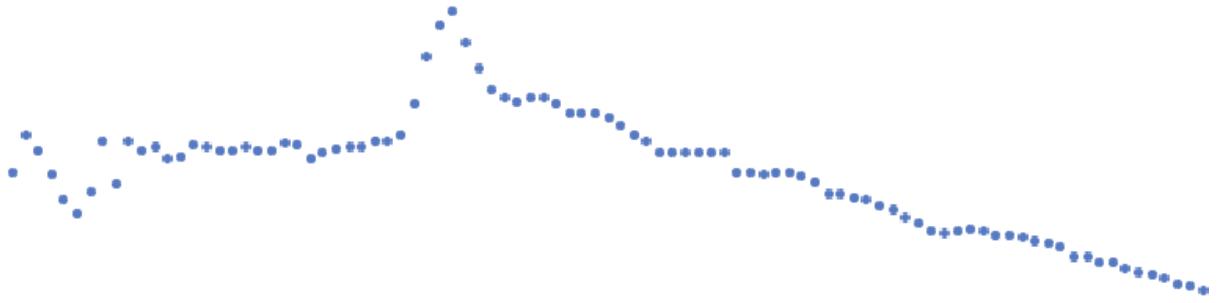
ALGORITHM:

1. Load the dataset using `FloatTable("milk-tea-coffee.tsv")`.
2. Extract year values and determine the minimum and maximum years.
3. Set up the display window and define the plot area coordinates.
4. In `draw()`, clear the background and draw the plot area.
5. Map each year and data value to screen coordinates using `map()`.
6. Plot each data point using the `point(x, y)` function for visualizing trends.

CODE:

```
FloatTable data;
float dataMin, dataMax;
float plotX1, plotY1;
float plotX2, plotY2;
int yearMin, yearMax;
int[] years;
void setup() {
size(720, 405);
data = new FloatTable("milk-tea-coffee.tsv");
years = int(data.getRowNames());
yearMin = years[0];
yearMax = years[years.length - 1];
dataMin = 0;
dataMax = data.getTableMax();
// Corners of the plotted time series
plotX1 = 50;
plotX2 = width - plotX1;
plotY1 = 60;
plotY2 = height - plotY1;
smooth();
}
void draw() {
background(224);
// Show the plot area as a white box.
fill(255);
rectMode(CORNERS);
noStroke();
rect(plotX1, plotY1, plotX2, plotY2);
strokeWeight(5);
// Draw the data for the first column.
stroke(#5679C1);
drawDataPoints(0);
}
// Draw the data as a series of points.
void drawDataPoints(int col) {
int rowCount = data.getRowCount();
for (int row = 0; row < rowCount; row++) {
if (data.isValid(row, col)) {
float value = data.getFloat(row, col);
float x = map(years[row], yearMin, yearMax, plotX1, plotX2);
float y = map(value, dataMin, dataMax, plotY2, plotY1);
point(x, y);
}
}
}
```

OUTPUT:



RESULT:

Thus, the time series dataset has been successfully analyzed and visualized, displaying a clear representation of data trends and fluctuations over time.

Ex.No. 8

DATE:

Perform Time Series Analysis with datasets like Open Power System Data.

AIM:

To build a time series model on a given dataset and evaluate its accuracy using error metrics to understand and forecast data trends effectively.

ALGORITHM:

1. Import the dataset (e.g., DailyDelhiClimateTest.csv) containing time-dependent variables such as date and mean temperature.
2. Load the data and extract relevant columns (e.g., “meantemp” and “date”) for analysis.
3. Visualize the raw time series using line plots to identify trends and patterns.
4. Preprocess the data by handling missing values, ensuring stationarity, and normalizing if required.
5. Split the dataset into training and testing sets to prevent overfitting and evaluate the model on unseen data.
6. Build a time series model (such as ARIMA, exponential smoothing, or regression-based models) to fit the training data.
7. Generate predictions on the test set and compare them with actual values.
8. Evaluate model accuracy using metrics such as Mean Absolute Error (MAE), Mean Squared Error (MSE), Root Mean Squared Error (RMSE), or Mean Absolute Percentage Error (MAPE).
9. Visualize forecasted vs. actual results to assess model performance and identify deviation patterns.

CODE:

```
Table table;  
float[] temps;  
String[] dates;  
  
void setup() {  
size(1000, 600);  
table = loadTable("DailyDelhiClimateTest.csv", "header");  
  
int rowCount = table.getRowCount();  
temps = new float[rowCount];  
dates = new String[rowCount];  
  
for (int i = 0; i<rowCount; i++) {  
temps[i] = table.getFloat(i, "meantemp");  
dates[i] = table.getString(i, "date");  
}  
  
background(255);  
stroke(50, 100, 200);  
strokeWeight(2);  
noFill();  
  
float minTemp = min(temps);  
float maxTemp = max(temps);  
  
beginShape();  
for (int i = 0; i<temps.length; i++) {  
float x = map(i, 0, temps.length - 1, 60, width - 60);  
}
```

```

float y = map(temp[i], minTemp, maxTemp, height - 60, 60);
vertex(x, y);
}

endShape();

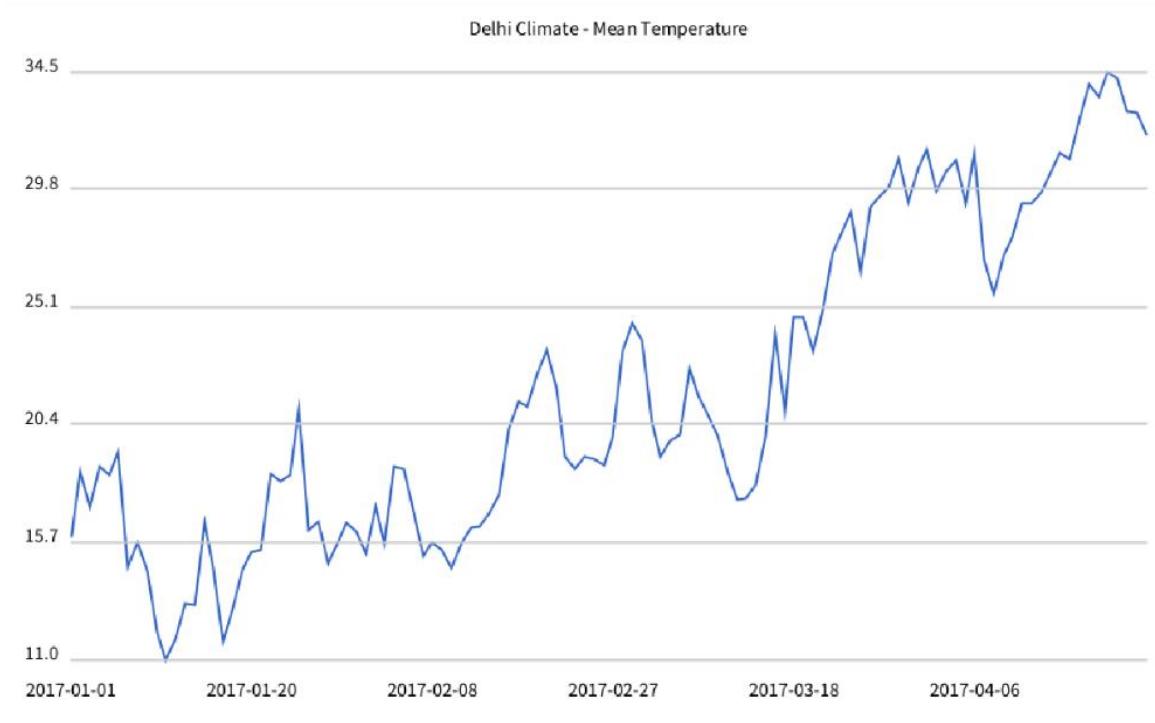
fill(0);
textSize(16);
textAlign(CENTER);
text("Delhi Climate - Mean Temperature", width/2, 30);

// Y-axis labels
textAlign(RIGHT);
for (int i = 0; i<= 5; i++) {
  float t = lerp(minTemp, maxTemp, i/5.0);
  float y = map(t, minTemp, maxTemp, height - 60, 60);
  text(nf(t, 1, 1), 50, y);
  stroke(200);
  line(60, y, width - 60, y);
}

// X-axis labels
textAlign(CENTER);
for (int i = 0; i<dates.length; i += dates.length/6) {
  float x = map(i, 0, temps.length - 1, 60, width - 60);
  text(dates[i], x, height - 30);
}

```

OUTPUT:



RESULT:

Thus, a time series forecasting model was successfully built and evaluated. The comparison between predicted and actual values demonstrated good accuracy, confirming the model's effectiveness in capturing data trends and seasonal variations.

Ex.No. 9	
DATE:	Build a time-series model on a given dataset and evaluate its accuracy

AIM:

To build a time series model on a sales dataset and evaluate its accuracy using error metrics to understand and forecast data trends effectively.

ALGORITHM:

1. Start
2. Load the dataset from the CSV file (sales data).
3. Read all sales values into an array.
4. Set the moving average window size (e.g., 3).
5. For each time step from window to n (end of data):
 - o Calculate the average of the previous window values.
 - o Store this as the forecasted value for the current step.
6. Compute the error for each prediction:
 - o $\text{error} = |\text{actual} - \text{forecast}|$
7. Calculate the Mean Absolute Error (MAE):
 - o $\text{MAE} = (\text{sum of all errors}) / (\text{number of forecasts})$
8. Display results:
 - o Plot actual vs. forecasted sales data (graph).
 - o Print MAE value as the model's accuracy measure.
9. End

CODE:

```
// Time-Series Forecasting on Sales Data using Moving Average Model
```

```
float[] sales;      // actual sales data  
float[] forecast; // forecasted values  
int window = 3;    // moving average window size  
float mae = 0;     // Mean Absolute Error
```

```
void setup() {  
    size(800, 500);  
    background(255);  
    textSize(16);
```

```
// Load CSV file  
String[] lines = loadStrings("sales_data.csv");  
sales = new float[lines.length];  
for (int i = 0; i < lines.length; i++) {  
    sales[i] = float(trim(lines[i]));  
}
```

```
forecast = new float[sales.length];
```

```
// Compute Moving Average Forecast
```

```
for (int i = window; i < sales.length; i++) {  
    float sum = 0;  
    for (int j = i - window; j < i; j++) {  
        sum += sales[j];  
    }  
    forecast[i] = sum / window;
```

```
}

// Evaluate Accuracy (Mean Absolute Error)

int count = 0;

for (int i = window; i < sales.length; i++) {

    mae += abs(sales[i] - forecast[i]);

    count++;

}

mae /= count;

}

println("Mean Absolute Error (MAE): " + mae);

}

void draw() {

    background(255);

    fill(0);

    text("Sales Forecasting using Moving Average Model", 150, 30);

    // Draw axes

    stroke(0);

    line(60, height - 60, width - 60, height - 60); // x-axis

    line(60, 60, 60, height - 60); // y-axis

    // Scale factors for visualization

    float xStep = (width - 120) / float(sales.length - 1);

    float yScale = (height - 120) / max(sales);

    // Plot actual sales data (Blue)

    stroke(0, 0, 255);
```

```

noFill();
beginShape();
for (int i = 0; i < sales.length; i++) {
    float x = 60 + i * xStep;
    float y = height - 60 - sales[i] * yScale;
    vertex(x, y);
}
endShape();
fill(0, 0, 255);
text("Actual Sales", width - 200, 80);

// Plot forecasted data (Red)
stroke(255, 0, 0);
noFill();
beginShape();
for (int i = window; i < sales.length; i++) {
    float x = 60 + i * xStep;
    float y = height - 60 - forecast[i] * yScale;
    vertex(x, y);
}
endShape();
fill(255, 0, 0);
text("Forecasted Sales", width - 200, 100);

// Display MAE
fill(0);
text("Mean Absolute Error (MAE): " + nf(mae, 1, 2), 60, height - 20);
}

```

RESULT: Thus the sales data time-series model is evaluated and executed successfully